

# Lab 3: Image Classification with CNN using PyTorch

## Objectives

This laboratory session focuses on implementing a Convolutional Neural Network (CNN) for satellite image classification using PyTorch. Students will gain hands-on experience in designing, training, and optimizing CNN architectures while working with real-world satellite imagery. The lab covers essential concepts including convolutional layers, pooling operations, batch normalization, dropout, and early stopping, culminating in model export and evaluation.

## Dataset Information

We will be using the "EuroSAT" dataset, which contains 27,000 labeled satellite images across 10 different land use and land cover classes. Each image is a 64x64 pixel RGB image. We use some of them classes for this lab.

Dataset link: <https://github.com/phelber/eurosat>

### Classes:

1. Annual Crop
2. Permanent Crop
3. River
4. Sea & Lake
5. Highway
6. Forest

## Tasks

### 1. Data Preparation

- Download and explore the EuroSAT dataset
- Implement data loading and preprocessing
- Create train/validation/test splits
- Apply data augmentation techniques

### 2. CNN Architecture Design

- Implement a basic CNN architecture
- Understand and configure kernel sizes, stride, and padding
- Add pooling layers (max and average pooling)
- Implement fully connected layers

### 3. Model Optimization

- Add Batch Normalization layers
- Implement Dropout for regularization
- Configure Early Stopping mechanism
- Experiment with different optimizers and learning rates

### 4. Training and Evaluation

- Train the model with and without optimization techniques
- Monitor training metrics
- Implement learning rate scheduling
- Evaluate model performance

## How CNN works?

A Convolutional Neural Network (CNN) works by processing an image through a series of layers to identify and classify the objects in it:

#### 1. Convolutional layer

- The first layer in a CNN, which uses filters to convolve with the image to create an activation map.

#### 2. Activation layer

- Applies a non-linear activation function, such as ReLU, to the output of the pooling layer. This helps the CNN learn more complex representations of the data.

#### 3. Pooling layer

- Uses filters to identify different parts of the image, such as edges and corners. The pooling layer reduces the dimensionality of the feature map, which speeds up computation and reduces memory.

#### 4. Fully connected layer

- Connects every neuron from the previous layer to every neuron in the fully connected layer. This layer integrates the features extracted in the previous layers and maps them to specific classes or outcomes.

```
In [55]: # !nvidia-smi
```

## Required Libraries and Setup

```
In [24]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
from torch.nn import functional as F
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, Dataset
import matplotlib.pyplot as plt
import numpy as np
from tqdm import tqdm
import os
```

```
In [25]: SEED = 1234

# random.seed(SEED)
# np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True
```

## Data Loading, Preprocessing and data loading

```
In [26]: image_folder = "D:\\Nokia_DL_L3_lab\\EuroSAT"
```

```
In [27]: mean = [0.5, 0.5, 0.5]
std = [0.5, 0.5, 0.5]
```

```
In [28]: def load_dataset(data_path):

    # Load all the images
    # Randomly augment the image data
    transformation = transforms.Compose([
        # Random horizontal flip
        transforms.RandomHorizontalFlip(0.5),
        # Random vertical flip
        transforms.RandomVerticalFlip(0.3),
        # transform to tensors
        transforms.ToTensor(),
        # Normalize the pixel values (in R, G, and B channels)
        transforms.Normalize(mean, std)
    ])

    # Load all of the images, transforming them
    full_dataset = torchvision.datasets.ImageFolder(
        root=data_path,
        transform=transformation
    )

    return full_dataset
```

```
In [29]: # Get the iterative dataloaders for test and training data
full_dataset = load_dataset(image_folder)
print("Data loaders ready to read", image_folder)
```

Data loaders ready to read D:\\Nokia\_DL\_L3\_lab\\EuroSAT

```
In [58]: classes = full_dataset.classes
```

```
In [59]: # Split into training (70% and testing (30%) datasets)
train_size = int(0.7 * len(full_dataset))
test_size = len(full_dataset) - train_size

# use torch.utils.data.random_split for training/test split
train_dataset, test_dataset = torch.utils.data.random_split(full_dataset, [train_size, test_size])

# define a loader for the training data we can iterate through in 50-image batches
train_loader = torch.utils.data.DataLoader(
    train_dataset,
    batch_size=500,
    num_workers=0,
    shuffle=False
)


# define a loader for the testing data we can iterate through in 50-image batches
test_loader = torch.utils.data.DataLoader(
    test_dataset,
```

```
batch_size=500,  
num_workers=0,  
shuffle=False  
)
```

```
In [60]: def imshow(inp):  
inp = inp.numpy().transpose((1, 2, 0))  
inp = std * inp + mean  
inp = np.clip(inp, 0, 1)  
plt.imshow(inp)  
plt.show()
```

## Sample Data

```
In [62]: imshow(train_dataset[0][0])  
print("Label: ", classes[train_dataset[0][1]])
```



Label: Forest

## CNN Model Architecture

```
In [35]: # Create a neural net class  
class CNNnet(nn.Module):  
  
    # Defining the Constructor  
    def __init__(self, num_classes=10):  
        super(CNNnet, self).__init__()  
  
        # In the init function, we define each layer we will use in our model  
  
        # Our images are RGB, so we have input channels = 3.  
        # We will apply 12 filters in the first convolutional layer  
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=20, kernel_size=3, stride=1, padding=1)  
  
        # A second convolutional layer takes 12 input channels, and generates 24 outputs  
        self.conv2 = nn.Conv2d(in_channels=20, out_channels=24, kernel_size=3, stride=1, padding=1)  
  
        # We in the end apply max pooling with a kernel size of 2  
        self.pool = nn.MaxPool2d(kernel_size=2)  
  
        # Our 64X64 image tensors will be pooled twice with a kernel size of 2. 64/2/2 is 16.  
        # This means that our feature tensors are now 16 x 16, and we've generated 24 of them  
  
        # We need to flatten these in order to feed them to a fully-connected layer  
        self.fc = nn.Linear(in_features=16 * 16 * 24, out_features=num_classes)  
  
    def forward(self, x):  
        # In the forward function, pass the data through the layers we defined in the init function  
  
        # Use a ReLU activation function after Layer 1 (convolution 1 and pool)  
        x = self.pool(F.relu(self.conv1(x)))  
  
        # Use a ReLU activation function after Layer 2  
        x = self.pool(F.relu(self.conv2(x)))  
  
        # Flatten  
        x = x.view(-1, 16 * 16 * 24)  
        # Feed to fully-connected layer to predict class  
        x = self.fc(x)
```

```
# Return class probabilities via a log_softmax function
return torch.log_softmax(x, dim=1)
```

```
In [36]: device = "cuda" if torch.cuda.is_available() else "cpu"
print(f"Using {device} device")
```

Using cuda device

```
In [37]: # Create an instance of the model class and allocate it to the device
model = CNNnet(num_classes=len(classes)).to(device)

print(model)
```

```
CNNnet(
  (conv1): Conv2d(3, 20, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(20, 24, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc): Linear(in_features=6144, out_features=6, bias=True)
)
```

## Training and Testing loops

```
In [63]: def train(model, device, train_loader, optimizer, epoch):

    # Set the model to training mode
    model.train()
    train_loss = 0
    print("Epoch:", epoch)

    # Process the images in batches
    for batch_idx, (data, target) in enumerate(train_loader):
        # Use the CPU or GPU as appropriate
        # Recall that GPU is optimized for the operations we are dealing with
        data, target = data.to(device), target.to(device)

        # Reset the optimizer
        optimizer.zero_grad()

        # Push the data forward through the model layers
        output = model(data)

        # print(output.shape)
        # print(target.shape)
        # Get the loss
        loss = loss_criteria(output, target)

        # Keep a running total
        train_loss += loss.item()

        # Backpropagate
        loss.backward()
        optimizer.step()

        # Print metrics so we see some progress
        # print('\tTraining batch {} Loss: {:.6f}'.format(batch_idx + 1, loss.item()))

    # return average loss for the epoch
    avg_loss = train_loss / (batch_idx+1)
    print('Training set: Average loss: {:.6f}'.format(avg_loss))
    return avg_loss
```

```
In [65]: def test(model, device, test_loader):
    # Switch the model to evaluation mode (so we don't backpropagate or drop)
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        batch_count = 0
        for data, target in test_loader:
            batch_count += 1
            data, target = data.to(device), target.to(device)

            # Get the predicted classes for this batch
            output = model(data)

            # Calculate the loss for this batch
            test_loss += loss_criteria(output, target).item()

            # Calculate the accuracy for this batch
            _, predicted = torch.max(output.data, 1)
            correct += torch.sum(target==predicted).item()

    # Calculate the average loss and total accuracy for this epoch
    avg_loss = test_loss / batch_count
    accuracy = 100. * correct / len(test_loader.dataset)
```

```
print('Testing set: Average loss: {:.6f}, Accuracy: {}/{} ({:.0f}%)\\n'.format(
    avg_loss, correct, len(test_loader.dataset),
    100. * correct / len(test_loader.dataset)))

# return average loss for the epoch
return avg_loss, accuracy
```

```
In [66]: # Use an "Adam" optimizer to adjust weights
optimizer = optim.Adam(model.parameters(), lr=0.01)

# Specify the Loss criteria
loss_criteria = nn.CrossEntropyLoss()

# Track metrics in these arrays
epoch_nums = []
training_loss = []
validation_loss = []

epochs = 20
print('Training on', device)
for epoch in range(1, epochs + 1):
    train_loss = train(model, device, train_loader, optimizer, epoch)
    test_loss, acc = test(model, device, test_loader)
    epoch_nums.append(epoch)
    training_loss.append(train_loss)
    validation_loss.append(test_loss)
```

Training on cuda  
Epoch: 1  
Training set: Average loss: 1.486689  
Testing set: Average loss: 0.783658, Accuracy: 3465/5100 (68%)

Epoch: 2  
Training set: Average loss: 0.664215  
Testing set: Average loss: 0.538084, Accuracy: 4041/5100 (79%)

Epoch: 3  
Training set: Average loss: 0.472326  
Testing set: Average loss: 0.504485, Accuracy: 4125/5100 (81%)

Epoch: 4  
Training set: Average loss: 0.409407  
Testing set: Average loss: 0.382344, Accuracy: 4387/5100 (86%)

Epoch: 5  
Training set: Average loss: 0.351420  
Testing set: Average loss: 0.358304, Accuracy: 4451/5100 (87%)

Epoch: 6  
Training set: Average loss: 0.316310  
Testing set: Average loss: 0.332839, Accuracy: 4505/5100 (88%)

Epoch: 7  
Training set: Average loss: 0.296192  
Testing set: Average loss: 0.384733, Accuracy: 4386/5100 (86%)

Epoch: 8  
Training set: Average loss: 0.303093  
Testing set: Average loss: 0.324226, Accuracy: 4513/5100 (88%)

Epoch: 9  
Training set: Average loss: 0.276628  
Testing set: Average loss: 0.321045, Accuracy: 4547/5100 (89%)

Epoch: 10  
Training set: Average loss: 0.274254  
Testing set: Average loss: 0.381354, Accuracy: 4404/5100 (86%)

Epoch: 11  
Training set: Average loss: 0.274323  
Testing set: Average loss: 0.327105, Accuracy: 4488/5100 (88%)

Epoch: 12  
Training set: Average loss: 0.252325  
Testing set: Average loss: 0.310457, Accuracy: 4535/5100 (89%)

Epoch: 13  
Training set: Average loss: 0.243192  
Testing set: Average loss: 0.313521, Accuracy: 4548/5100 (89%)

Epoch: 14  
Training set: Average loss: 0.230995  
Testing set: Average loss: 0.302563, Accuracy: 4569/5100 (90%)

Epoch: 15  
Training set: Average loss: 0.226927  
Testing set: Average loss: 0.292644, Accuracy: 4572/5100 (90%)

Epoch: 16  
Training set: Average loss: 0.246186  
Testing set: Average loss: 0.287697, Accuracy: 4599/5100 (90%)

Epoch: 17  
Training set: Average loss: 0.219593  
Testing set: Average loss: 0.289952, Accuracy: 4602/5100 (90%)

Epoch: 18  
Training set: Average loss: 0.209093  
Testing set: Average loss: 0.263947, Accuracy: 4626/5100 (91%)


Epoch: 19  
Training set: Average loss: 0.194871  
Testing set: Average loss: 0.290088, Accuracy: 4596/5100 (90%)

Epoch: 20  
Training set: Average loss: 0.205620  
Testing set: Average loss: 0.265503, Accuracy: 4638/5100 (91%)

## Plotting training and validation losses

```
In [67]: plt.figure()  
         plt.plot(epoch_nums, training_loss)
```

```
plt.plot(epoch_nums, validation_loss)
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['training', 'validation'], loc='upper right')
plt.show()
```



## Training Loop with Optimization Techniques - With batch normalization and early stopping

```
In [68]: class CNNnet1(nn.Module):
def __init__(self, num_classes=10):
    super(CNNnet1, self).__init__()

    # First convolutional block
    self.conv1 = nn.Conv2d(in_channels=3, out_channels=20, kernel_size=3, stride=1, padding=1)
    # self.conv2 = nn.Conv2d(in_channels=3, out_channels=20, kernel_size=3, stride=1, padding=1)
    self.batch_norm1 = nn.BatchNorm2d(20)

    # Second convolutional block
    self.conv2 = nn.Conv2d(in_channels=20, out_channels=24, kernel_size=3, stride=1, padding=1)
    self.batch_norm2 = nn.BatchNorm2d(24)

    # Pooling and dropout
    self.pool = nn.MaxPool2d(kernel_size=2)
    self.drop = nn.Dropout2d(p=0.2)

    # Fully connected layer
    self.fc = nn.Linear(in_features=16 * 16 * 24, out_features=num_classes)

def forward(self, x):
    # First block: Conv -> BatchNorm -> ReLU -> Pool
    x = self.conv1(x)
    # x = self.batch_norm1(x)
    x = F.relu(x)
    x = self.pool(x)

    # Second block: Conv -> BatchNorm -> ReLU -> Pool
    x = self.conv2(x)
    x = self.batch_norm2(x)
    x = F.relu(x)
    x = self.pool(x)

    # Dropout
    x = self.drop(x)

    # Flatten and feed to fully connected layer
    x = x.view(-1, 16 * 16 * 24)
    x = self.fc(x)

    return torch.log_softmax(x, dim=1)
```

```
In [69]: # Use an "Adam" optimizer to adjust weights
model1 = CNNnet1(len(classes)).to(device)
optimizer1 = optim.Adam(model1.parameters(), lr=0.01)

# Specify the Loss criteria
loss_criteria1 = nn.CrossEntropyLoss()

# Track metrics in these arrays
epoch_nums1 = []
training_loss1 = []
```

```
validation_loss1 = []
```

```
print(model1)
```

```
CNNnet1(  
    (conv1): Conv2d(3, 20, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (batch_norm1): BatchNorm2d(20, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (conv2): Conv2d(20, 24, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (batch_norm2): BatchNorm2d(24, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (drop): Dropout2d(p=0.2, inplace=False)  
    (fc): Linear(in_features=6144, out_features=6, bias=True)  
)
```

```
In [70]: epochs = 30
```

```
best_accuracy = -1
```

```
early_stop_thresh = 5
```

```
print('Training on', device)
```

```
for epoch in range(1, epochs + 1):  
    train_loss = train(model1, device, train_loader, optimizer1, epoch)  
    test_loss, acc = test(model1, device, test_loader)  
  
    if acc > best_accuracy:  
        best_accuracy = acc  
        best_epoch = epoch  
    elif epoch - best_epoch > early_stop_thresh:  
        print("Early stopped training at epoch %d" % epoch)  
        break # terminate the training loop  
  
    epoch_nums1.append(epoch)  
    training_loss1.append(train_loss)  
    validation_loss1.append(test_loss)
```



Training on cuda  
Epoch: 1  
Training set: Average loss: 6.051547  
Testing set: Average loss: 2.471620, Accuracy: 1855/5100 (36%)

Epoch: 2  
Training set: Average loss: 1.728463  
Testing set: Average loss: 1.308100, Accuracy: 2590/5100 (51%)

Epoch: 3  
Training set: Average loss: 1.316040  
Testing set: Average loss: 1.305446, Accuracy: 2750/5100 (54%)

Epoch: 4  
Training set: Average loss: 1.191673  
Testing set: Average loss: 1.754185, Accuracy: 2381/5100 (47%)

Epoch: 5  
Training set: Average loss: 1.036223  
Testing set: Average loss: 1.000470, Accuracy: 2967/5100 (58%)

Epoch: 6  
Training set: Average loss: 0.855498  
Testing set: Average loss: 0.843346, Accuracy: 3475/5100 (68%)

Epoch: 7  
Training set: Average loss: 0.704369  
Testing set: Average loss: 0.704968, Accuracy: 3912/5100 (77%)

Epoch: 8  
Training set: Average loss: 0.604896  
Testing set: Average loss: 0.867081, Accuracy: 3541/5100 (69%)

Epoch: 9  
Training set: Average loss: 0.564415  
Testing set: Average loss: 0.592206, Accuracy: 4093/5100 (80%)

Epoch: 10  
Training set: Average loss: 0.532608  
Testing set: Average loss: 0.452310, Accuracy: 4263/5100 (84%)

Epoch: 11  
Training set: Average loss: 0.497118  
Testing set: Average loss: 0.449877, Accuracy: 4301/5100 (84%)

Epoch: 12  
Training set: Average loss: 0.485523  
Testing set: Average loss: 0.561836, Accuracy: 3907/5100 (77%)

Epoch: 13  
Training set: Average loss: 0.464917  
Testing set: Average loss: 0.450655, Accuracy: 4277/5100 (84%)

Epoch: 14  
Training set: Average loss: 0.464591  
Testing set: Average loss: 1.591951, Accuracy: 3467/5100 (68%)

Epoch: 15  
Training set: Average loss: 0.450793  
Testing set: Average loss: 0.449253, Accuracy: 4243/5100 (83%)

Epoch: 16  
Training set: Average loss: 0.370977  
Testing set: Average loss: 0.356077, Accuracy: 4463/5100 (88%)

Epoch: 17  
Training set: Average loss: 0.397162  
Testing set: Average loss: 0.409760, Accuracy: 4349/5100 (85%)

Epoch: 18  
Training set: Average loss: 0.377337  
Testing set: Average loss: 0.383617, Accuracy: 4325/5100 (85%)

Epoch: 19  
Training set: Average loss: 0.332562  
Testing set: Average loss: 0.428558, Accuracy: 4246/5100 (83%)

Epoch: 20  
Training set: Average loss: 0.319862  
Testing set: Average loss: 0.308353, Accuracy: 4528/5100 (89%)

Epoch: 21  
Training set: Average loss: 0.308724  
Testing set: Average loss: 0.851021, Accuracy: 3756/5100 (74%)

Epoch: 22  
Training set: Average loss: 0.341983

Testing set: Average loss: 0.321126, Accuracy: 4509/5100 (88%)

Epoch: 23

Training set: Average loss: 0.289640

Testing set: Average loss: 0.312582, Accuracy: 4526/5100 (89%)

Epoch: 24

Training set: Average loss: 0.302459

Testing set: Average loss: 0.434318, Accuracy: 4233/5100 (83%)

Epoch: 25

Training set: Average loss: 0.288352

Testing set: Average loss: 0.348134, Accuracy: 4434/5100 (87%)


Epoch: 26

Training set: Average loss: 0.269218

Testing set: Average loss: 0.317410, Accuracy: 4491/5100 (88%)

Early stopped training at epoch 26

```
In [46]: plt.figure()
plt.plot(epoch_nums1, training_loss1)
plt.plot(epoch_nums1, validation_loss1)
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['training', 'validation'], loc='upper right')
plt.show()
```




## Testing on random images

```
In [47]: def predict_img(img, model):
# data, target = data.to(device), target.to(device)
xb = img.unsqueeze(0).to(device)
yb = model(xb)

_, pred = torch.max(yb, dim=1)
return pred
```

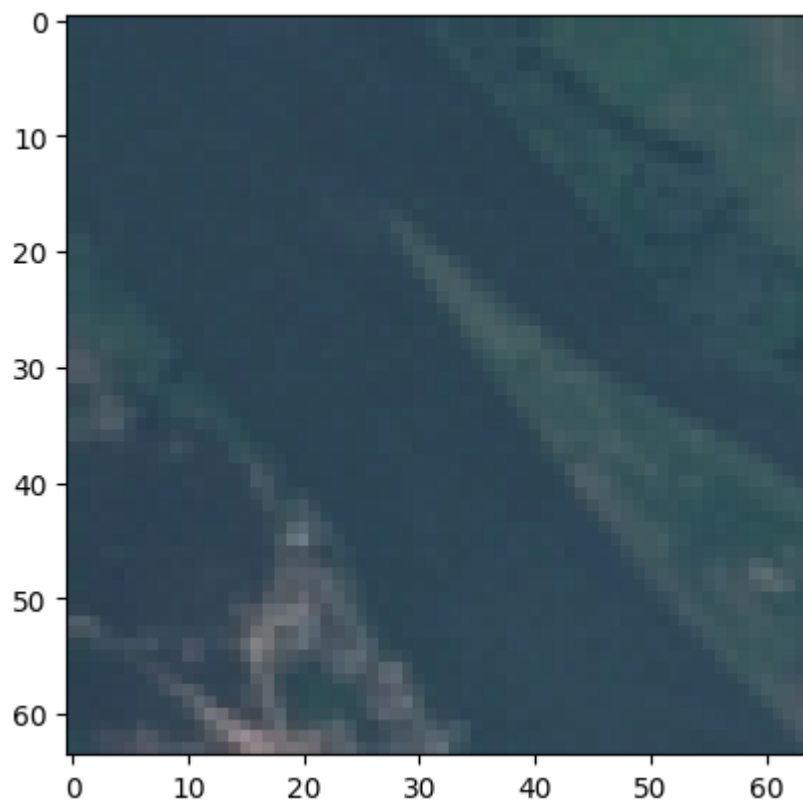
```
In [49]: img, label = test_dataset[10]
imshow(img)
print('Label:', classes[label], ', Predicted:', classes[predict_img(img, model1).item()])
```



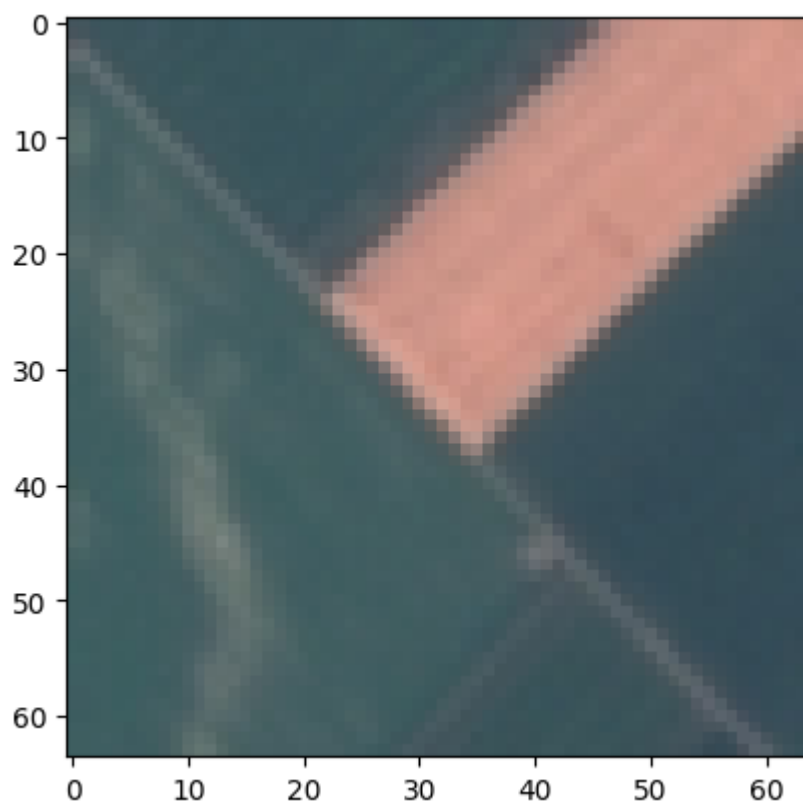
Label: River , Predicted: River

```
In [55]: for i in torch.randint(low=1, high=5000, size=(5,)):
# print(i.item())
img, label = test_dataset[i.item()]
print('Label:', classes[label], ', Predicted:', classes[predict_img(img, model1).item()])
imshow(img)
```

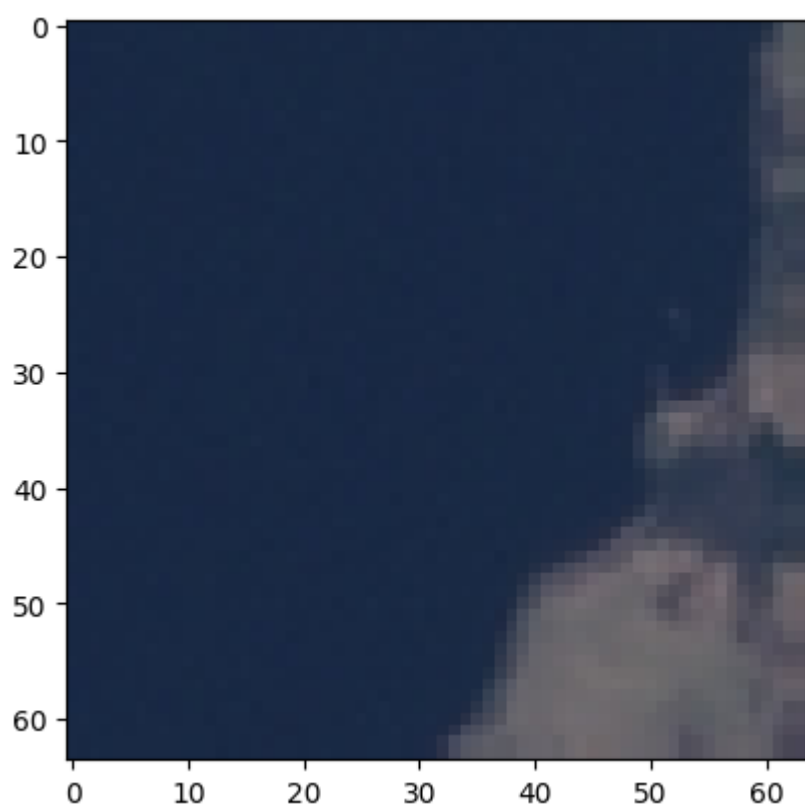
Label: River , Predicted: River




Label: AnnualCrop , Predicted: AnnualCrop



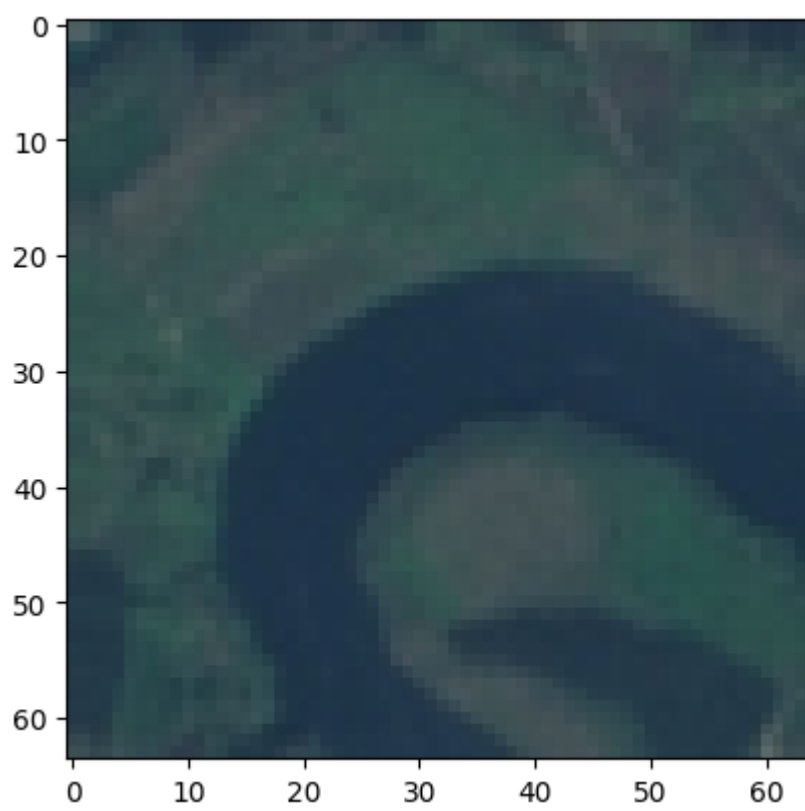
Label: SeaLake , Predicted: SeaLake



Label: Highway , Predicted: Highway



Label: River , Predicted: SeaLake



In [ ]: