

Lab 7: Implementing Object Detection with YOLO for Telecom Infrastructure

Objective This lab focuses on implementing object detection for telecom infrastructure using YOLOv8. Students will learn to train and deploy a YOLO model for detecting different types of antennas (GSM and Microwave) in real-world scenarios. The lab demonstrates practical applications of object detection in the telecommunications industry, covering model training, evaluation, and inference on new images.

Dataset Information

Source: [Tower Detection Dataset](#)

Classes: 3 types of antennas

- GSM Antenna
- GSM Antenna (variant)
- Microwave Antenna

Dataset Split

- Training: 1350 images (99%)
- Validation: 7 images (1%)
- Testing: 5 images

Tasks Overview

- Environment Setup and Dependencies Installation
- Dataset Download and Preparation
- YOLOv8 Model Configuration
- Model Training
- Performance Evaluation
- Inference on Test Images
- Model Export

Task 1: Environment Setup and Dependencies Installation

```
# Install required packages
!pip install ultralytics
!pip install roboflow

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.1.4->ultralytics) (2024.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (3.10)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (2.2)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (202)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.16.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (4)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.4.2)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.1.4)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (2024.10.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch>=1.8.0->ultra)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib>=3.3.0->ultr)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.8.0->ultralytics) (3)
Downloading ultralytics-8.3.51-py3-none-any.whl (901 kB)
901.3/901.3 kB 20.1 MB/s eta 0:00:00
Downloading ultralytics_thop-2.0.13-py3-none-any.whl (26 kB)
Installing collected packages: ultralytics-thop, ultralytics
Successfully installed ultralytics-8.3.51 ultralytics-thop-2.0.13
Collecting roboflow
  Downloading roboflow-1.1.50-py3-none-any.whl.metadata (9.7 kB)
Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-packages (from roboflow) (2024.12.14)
Collecting idna==3.7 (from roboflow)
  Downloading idna-3.7-py3-none-any.whl.metadata (9.9 kB)
Requirement already satisfied: cyclor in /usr/local/lib/python3.10/dist-packages (from roboflow) (0.12.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.10/dist-packages (from roboflow) (1.4.7)
```

```
installing collected packages: filetype, python-dotenv, idna, roboflow
Attempting uninstall: idna
  Found existing installation: idna 3.10
  Uninstalling idna-3.10:
    Successfully uninstalled idna-3.10

# Import necessary libraries
import ultralytics
from ultralytics import YOLO
from roboflow import Roboflow
import os
import cv2
import numpy as np
from google.colab import drive

🔗 Creating new Ultralytics Settings v0.0.6 file ✅
View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see 

## Task 2: Dataset Download and Preparation


```

```
# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')

# Set working directory
import os
%cd /content

🔗 Mounted at /content/drive
/content
```

Setup Data path

```
# Define dataset path (update this to your Google Drive path)
DATASET_PATH = '/content/drive/MyDrive/tower_dataset' # Update this path

# Verify dataset structure
!ls {DATASET_PATH}

# Display data.yaml content
!cat {DATASET_PATH}/data.yaml

# Setup directory structure
train_path = os.path.join(DATASET_PATH, 'train/images')
valid_path = os.path.join(DATASET_PATH, 'valid/images')
test_path = os.path.join(DATASET_PATH, 'test/images')

# Verify image counts
print(f"Training images: {len(os.listdir(train_path))}")
print(f"Validation images: {len(os.listdir(valid_path))}")
print(f"Test images: {len(os.listdir(test_path))}")

🔗 best.pt data.yaml test train valid
train: ../train/images
val: ../valid/images
test: ../test/images

nc: 3
names: ['GSM Antenna', 'GSM Antenna ', 'Microwave Antenna']

roboflow:
  workspace: object-detection-yolo-c8gsd
  project: tower-detection-tff1p
  version: 3
  license: CC BY 4.0
  url: https://universe.roboflow.com/object-detection-yolo-c8gsd/tower-detection-tff1p/dataset/3Training images: 1350
Validation images: 7
Test images: 5
```

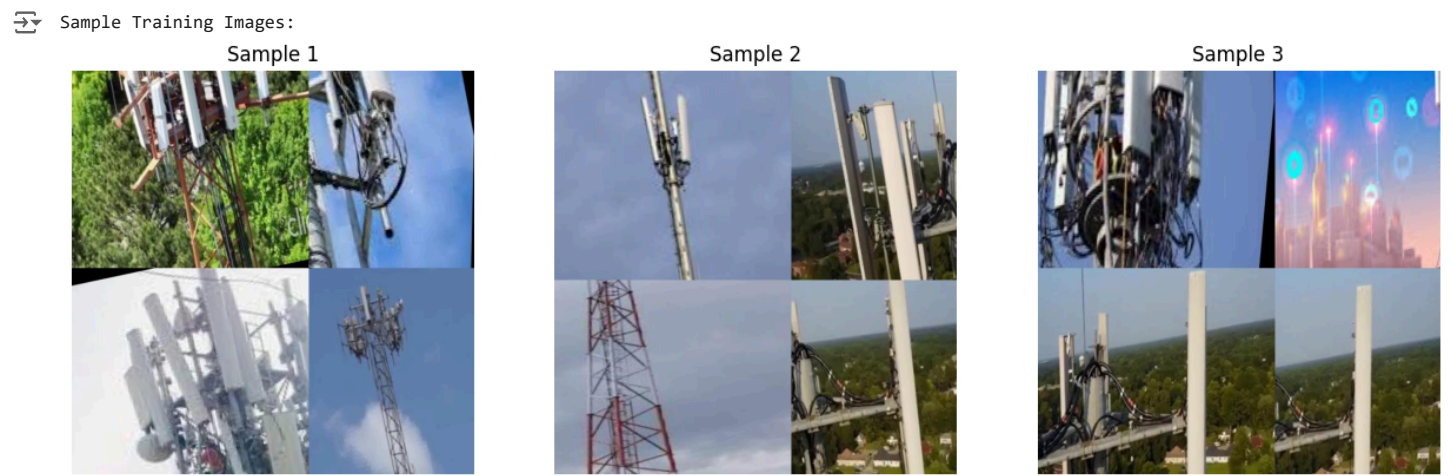
Data Visualization

```
import matplotlib.pyplot as plt
import random

def display_sample_images(image_path, num_samples=3):
    """Display random sample images from the dataset"""
    images = os.listdir(image_path)
    samples = random.sample(images, num_samples)

    plt.figure(figsize=(15, 5))
    for i, img_name in enumerate(samples, 1):
        img = cv2.imread(os.path.join(image_path, img_name))
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        plt.subplot(1, num_samples, i)
        plt.imshow(img)
        plt.title(f'Sample {i}')
        plt.axis('off')
    plt.show()

# Display sample training images
print("Sample Training Images:")
display_sample_images(train_path)
```



Task 3: YOLOv8 Model Configuration

```
import torch

# Check CUDA availability
print(f"CUDA is available: {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"CUDA Device: {torch.cuda.get_device_name(0)}")
    print(f"Number of CUDA devices: {torch.cuda.device_count()}")
    print(f"CUDA Version: {torch.version.cuda}")

🔗 CUDA is available: True
   CUDA Device: Tesla T4
   Number of CUDA devices: 1
   CUDA Version: 12.1

# Load a pre-trained YOLOv8 model
model = YOLO('yolov8n.pt')

# Move model to GPU if available
device = 'cuda' if torch.cuda.is_available() else 'cpu'
print(f"\nUsing device: {device}")

# Display model information
print("\nModel Information:")
print(model.info())

# Display model parameters
print(f"\nTotal Parameters: {sum(p.numel() for p in model.parameters())}")

🔗 Downloading https://github.com/ultralytics/assets/releases/download/v8.3.0/yolov8n.pt to 'yolov8n.pt'...
   100% ██████████ 6.25M/6.25M [00:00<00:00, 362MB/s]

Using device: cuda

Model Information:
YOLOv8n summary: 225 layers, 3,157,200 parameters, 0 gradients, 8.9 GFLOPs
(225, 3157200, 0, 8.8575488)

Total Parameters: 3157200
```

Task 4: Model Training

Configure Training Parameters

```
# Define training configuration
train_config = {
    'data': os.path.join(DATASET_PATH, 'data.yaml'),
    'epochs': 50,
    'imgsz': 640,
    'batch': 32,
    'name': 'tower_detection_model',
    'patience': 20, # Early stopping patience
    'save': True, # Save best model
    'device': 0 if torch.cuda.is_available() else 'cpu', # Use GPU if available
    'workers': 8, # Number of worker threads
    'optimizer': 'Adam', # Optimizer (SGD, Adam, AdamW)
    'lr0': 0.01, # Initial learning rate
    'weight_decay': 0.0005, # Weight decay
    'exist_ok': True, # Overwrite existing experiment
    'pretrained': True, # Use pretrained backbone
    'amp': True, # Automatic Mixed Precision
}


# Adjust batch size based on available GPU memory
if torch.cuda.is_available():
    gpu_mem = torch.cuda.get_device_properties(0).total_memory / 1e9 # Memory in GB
    print(f"\nGPU Memory Available: {gpu_mem:.2f} GB")

    # Adjust batch size based on GPU memory
    if gpu_mem < 8:
```

```

    train_config['batch'] = 8
elif gpu_mem < 16:
    train_config['batch'] = 16
else:
    train_config['batch'] = 32

print("\nTraining Configuration:")
for key, value in train_config.items():
    print(f"{key}: {value}")
```


 GPU Memory Available: 15.84 GB

Training Configuration:
data: /content/drive/MyDrive/tower_dataset/data.yaml
epochs: 50
imgsz: 640
batch: 16
name: tower_detection_model
patience: 20
save: True
device: 0
workers: 8
optimizer: Adam
lr0: 0.01
weight_decay: 0.0005
exist_ok: True
pretrained: True
amp: True

Start Training

Start training with configured parameters
results = model.train(**train_config)

Save training results
model.save(os.path.join(DATASET_PATH, 'best.pt'))



42/50	2.56G	1.679	1.476	1.648	76	640: 100%	<div></div>	85/85 [00:29<00:00, 2.87it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	<div></div> 1/1 [00:00<00:00, 3.34it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

43/50 2.54G 1.68 1.462 1.651 59 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 6.04it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

44/50 2.59G 1.672 1.448 1.65 69 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 6.46it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

45/50 2.58G 1.672 1.441 1.644 59 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 3.62it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

46/50 2.59G 1.638 1.406 1.616 71 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 6.04it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

47/50 2.59G 1.634 1.409 1.617 54 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 4.86it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

48/50 2.58G 1.61 1.388 1.606 54 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 5.73it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

49/50 2.61G 1.605 1.37 1.593 77 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 4.52it/s]

Epoch GPU_mem box_loss cls_loss dfl_loss Instances Size

50/50 2.57G 1.583 1.365 1.594 97 640: 100%

Class Images Instances Box(P R mAP50 mAP50-95): 100%

1/1 [00:00<00:00, 2.89it/s]

50 epochs completed in 0.503 hours.
Optimizer stripped from runs/detect/tower_detection_model/weights/last.pt, 6.2MB
Optimizer stripped from runs/detect/tower_detection_model/weights/best.pt, 6.2MB

Validating runs/detect/tower_detection_model/weights/best.pt...
Ultralytics 8.3.51 Python-3.10.12 torch-2.5.1+cu121 CUDA:0 (Tesla T4, 15102MiB)
Model summary (fused): 168 layers, 3,006,233 parameters, 0 gradients, 8.1 GFLOPs

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	<div></div>	1/1 [00:00<00:00, 8.02it/s]
all	7	21	0.749	0.656	0.75	0.402		
GSM Antenna	6	16	0.721	0.312	0.505	0.18		
Microwave Antenna	4	5	0.778	1	0.995	0.624		

Speed: 0.2ms preprocess, 3.5ms inference, 0.0ms loss, 4.2ms postprocess per image
Results saved to runs/detect/tower_detection_model

Task 5: Performance Evaluation

```
# Plot training metrics
from ultralytics.utils.plotting import plot_results
```

```
# Plot the results using the static method
plot_results(file='/content/runs/detect/tower_detection_model/results.png') # plot results.txt as results.png
plt.show()

# Display final metrics
print("\nTraining Results:")

# Load and display metrics from results.csv
import pandas as pd
try:
    results_df = pd.read_csv('/content/runs/detect/tower_detection_model/results.csv')
    print("\nFinal Metrics:")
    print(f"Best mAP50: {results_df['metrics/mAP50(B)'].max():.4f}")
    print(f"Best mAP50-95: {results_df['metrics/mAP50-95(B)'].max():.4f}")
    print(f"Final Epoch: {len(results_df)}")
except Exception as e:
    print(f"Could not load results.csv: {e}")
```

↗

Training Results:

Final Metrics:
Best mAP50: 0.7504
Best mAP50-95: 0.4021
Final Epoch: 50

Performance Evaluation

```
# Validate the model
metrics = model.val()

# Display metrics
print(f"mAP50: {metrics.box.map50}")
print(f"mAP50-95: {metrics.box.map}")

↗ Ultralytics 8.3.51 🚀 Python-3.10.12 torch-2.5.1+cu121 CUDA:0 (Tesla T4, 15102MiB)
Model summary (fused): 168 layers, 3,006,233 parameters, 0 gradients, 8.1 GFLOPs
val: Scanning /content/drive/MyDrive/tower_dataset/valid/labels.cache... 7 images, 0 backgrounds, 0 corrupt: 100%|██████████| 7/7 [00:00
Class  Images  Instances  Box(P      R      mAP50  mAP50-95): 100%|██████████| 1/1 [00:00<00:00,  2.90it/s]
    all             7         21    0.644    0.656      0.7      0.365
  GSM Antenna       6         16    0.658    0.312    0.505    0.179
Microwave Antenna   4          5    0.631      1     0.895    0.552
Speed: 0.3ms preprocess, 16.9ms inference, 0.0ms loss, 2.0ms postprocess per image
Results saved to runs/detect/tower_detection_model
mAP50: 0.699988939680277
mAP50-95: 0.36516746670420464
```

✓ Task 6: Inference on Test Images

```
def predict_image(image_path):
    # Perform prediction
    results = model.predict(image_path, conf=0.25)

    # Process and display results
    for result in results:
        boxes = result.boxes
        for box in boxes:
            # Get coordinates and class
            x1, y1, x2, y2 = box.xyxy[0]
            cls = box.cls
            conf = box.conf


            # print(f"Detected {model.names[int(cls)]} with confidence {conf:.2f}")

    return results[0].plot()

# Test on sample image
test_image = "/content/drive/MyDrive/tower_dataset/test/images/MP_3_jpeg.rf.7bded774087bd0b184d65eb93ce0eb04.jpg"
prediction = predict_image(test_image)

↗ image 1/1 /content/drive/MyDrive/tower_dataset/test/images/MP_3_jpeg.rf.7bded774087bd0b184d65eb93ce0eb04.jpg: 640x640 1 Microwave Antenn
Speed: 3.9ms preprocess, 12.7ms inference, 1.7ms postprocess per image at shape (1, 3, 640, 640)
```

```
#Plot the show image of prediction
plt.figure(figsize=(10, 10))
plt.imshow(prediction)
plt.axis('off')
plt.show()
```



Task 7: Model Export

```
# Export model to ONNX format
model.export(format='onnx')
```